Grantwork

Grant funding — a Flowgrid OS module for applications & impact reporting

AVAILABLE IN PILOT
Enabled per client workspace
flowgridos.co.uk

FOR CHARITIES, CICS AND THE CONSULTANTS WHO FUNDRAISE FOR THEM

Keep the application current. The funder's monitoring return assembles itself.

Winning the grant is the easy half. What follows is years of reporting against outcomes on someone else's schedule — today held in a spreadsheet, a drive of past bids, and one person's memory of what each funder wants. Grantwork holds the portfolio and the obligations, then drafts the returns from the outcomes you actually recorded.

Case for support → Prospect → Apply → Decision → Deliver → Monitor → Evaluate

What it keeps hold of

- **Funders and a seeded UK catalogue** — lottery funds, trusts, community foundations — each row carrying the date its criteria were last verified.
- **Applications** through the pipeline, with amounts requested and awarded, restricted or unrestricted.
- **Award conditions** — including which must be discharged before you can draw down.
- **Impact measures** with baseline, target and unit, and the outcomes recorded against each reporting period.
- **The obligation calendar** — every reporting period, its due date, and what is overdue.

Numbers a funder can trust

- **Every figure comes from a recorded outcome.** The impact table renders from your rows; the AI writes only the narrative around it.
- **An unmeasured outcome says "not recorded"** — never a zero, and never a plausible estimate. Absence must not read as failure, or as success.
- **Over-delivery is shown as achieved**, uncapped, rather than quietly rounded to target.
- **Unverified funder criteria warn on page one** of any bid built on them.
- **Draft-first with a DRAFT watermark.** Nothing reaches a funder without a human sending it.

What it drafts

Case for support
The anchor document, grounded in your own need evidence and beneficiary stories.

Funding application
Written to the funder's own criteria as recorded in the catalogue.

Monitoring return
Nine sections built from the period's recorded outcomes and conditions.

Impact evaluation
End of grant, aggregating every reporting period.

Impact card
One-page funder-facing PDF in your brand colour. No AI — pure calculation.

- ✓ **Success rate by funder**
  And pipeline value weighted by stage.
- ✓ **Restricted vs unrestricted split**
  Calculated, not tracked by hand.
- ✓ **Funder concentration risk**
  What share of income rests on your largest funder.
- ✓ **A tab on any project**
  Grant applications sit alongside a development scheme: a CLT building homes is also a grant-funded charity.

Why this matters commercially. The sector spends on the order of **£900m a year applying for grants**, and charities in the £100k–£1m band spend around **35% of their grant income** on the applications themselves. Every hour returned to delivery is the product.

9 / 9
sections of a live monitoring return drafted from recorded data in about a minute, impact table from real figures

0
invented citations in a live case for support — every source claim resolved to its page

0.7 sec
to render a funder-facing impact card, with no model call at all

Who it's for

- **Charities, CICs and community organisations** with roughly £100k–£3m income, running a rolling portfolio of applications and the reporting that follows.
- **Fundraising consultancies and freelance bid writers** — one branded workspace per client, each sealed off from the others at the database, carrying your name rather than ours.
- **Infrastructure bodies** — CVSs, community foundations and networks supporting many small organisations at once.

What it deliberately does not do

- **No donor fundraising.** Individual giving, direct debits and Gift Aid stay in your CRM — this is the institutional grant side only.
- **No restricted-fund accounting.** It records what was awarded and reported, not your ledger.
- **No funder portal or auto-submission.** Returns are drafted for a human to send, never filed for you.
- **No impact framework beyond measures against targets.** Deliberately simple, because the figures have to be defensible to a funder.

Narrow on purpose. One loop — application current, return assembles itself — done properly, rather than a second CRM.

Bring one live grant. See the return it writes.
We will set up one real application with its outcomes, and draft the next monitoring return with you.

Book a working session